


# Next '22

## Better Hardware Provisioning for ML Experiments on GCP

Provision hardware for ML reliably.

Oct/  
11-13



# Sayak Paul

## ML Engineer at Carted

October 13, 2022





I'm presenting work done by my  
colleagues (Nilabhra, Shane, Sam)  
and myself.


# Contents

- 01 Typical workflow for ML hardware provisioning

---
- 02 Problems with the typical workflow

---
- 03 Streamlining provisioning workflow with Terraform

---
- 04 Demo



# The standard workflow

# Components of ML experiments

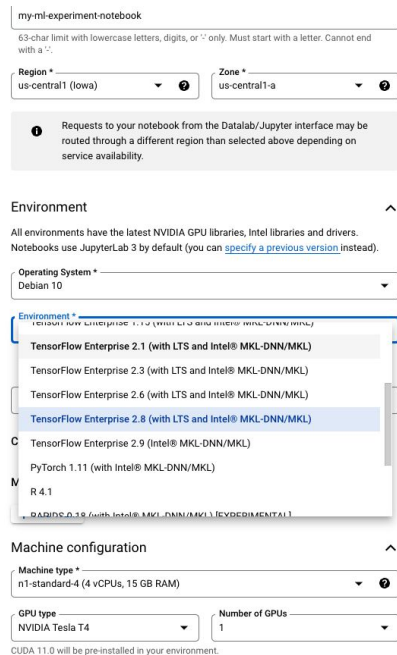
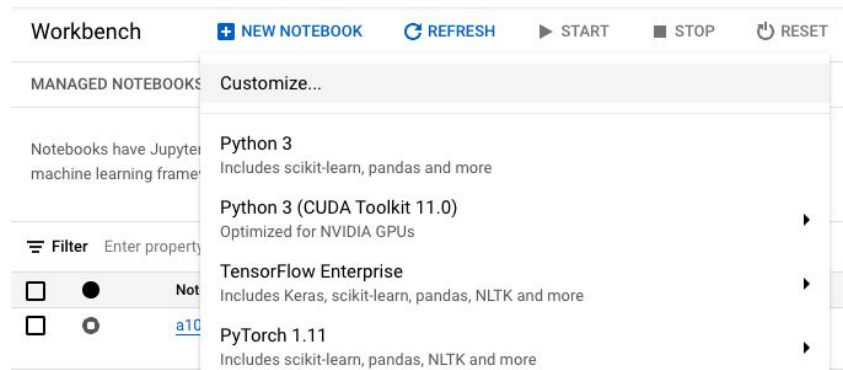
---

We need:

- Pre-configured environment (with libraries and CUDA stuff sorted).
- Notebook instance for better interaction with the modules.
- Access to a GPU or GPUs.
- Optionally a central storage location for storing experiment artifacts.

# Vertex AI and GCS to the rescue

Vertex AI Workbench allows to do this with a few clicks!



# Vertex AI and GCS to the rescue

Same applies to GCS too!



← Create a bucket

- Name your bucket**  
Pick a globally unique, permanent name. [Naming guidelines](#)  
  
Tip: Don't include any sensitive information  
✓ LABELS (OPTIONAL)
- Choose where to store your data**  
Location: us (multiple regions in United States)  
Location type: Multi-region
- Choose a default storage class for your data**  
Default storage class: Standard
- Choose how to control access to objects**  
Public access prevention: On  
Access control: Uniform
- Choose how to protect object data**  
Protection tools: None  
Data encryption: Google-managed key



# Problems with the typical workflow



# Handling hardware components


---

- Manual efforts
  - Redundancy from repeating the instructions every time a new instance needs to be created.
  - What if a step is missed out?
- Easy to forget about cleaning up the resources after experimentation
  - Multiple moving pieces (notebook instance, GCS bucket, etc.)



Provisioning infrastructure through point-and-click GUI's or custom scripts is slow, error-prone, inefficient, and doesn't scale."

Terraform



# Further extensions

---


- Supporting custom dependencies
- Supporting Python virtual environments

# Observations

---

- **Bad:**
  - Multiple interconnected steps that can be easily missed out and wreak havoc
- **Good:**
  - Single commands to provision and deprovision the resources

Enter  
Terraform



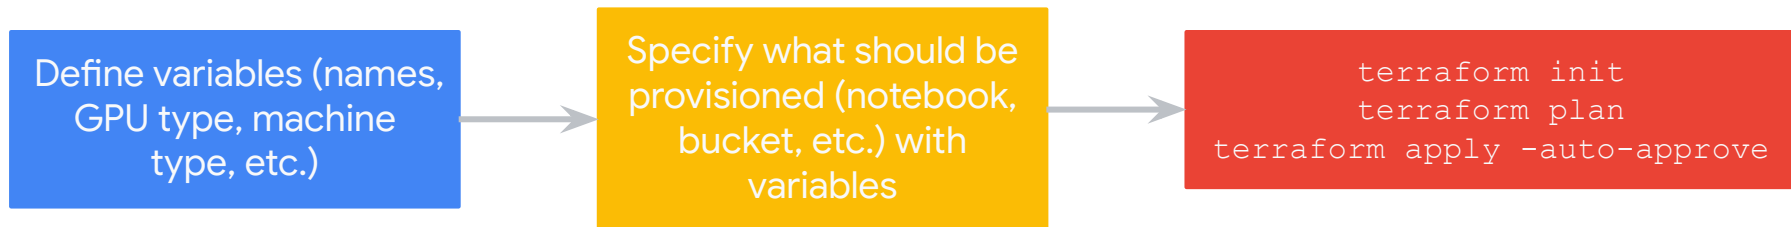
With HashiCorp Terraform, provisioning and security can be automated based on infrastructure and policy as code. Infrastructure and policies are codified, shared, managed, and executed within a workflow that is consistent across all infrastructure.

Terraform



# Realize the same workflow but with code

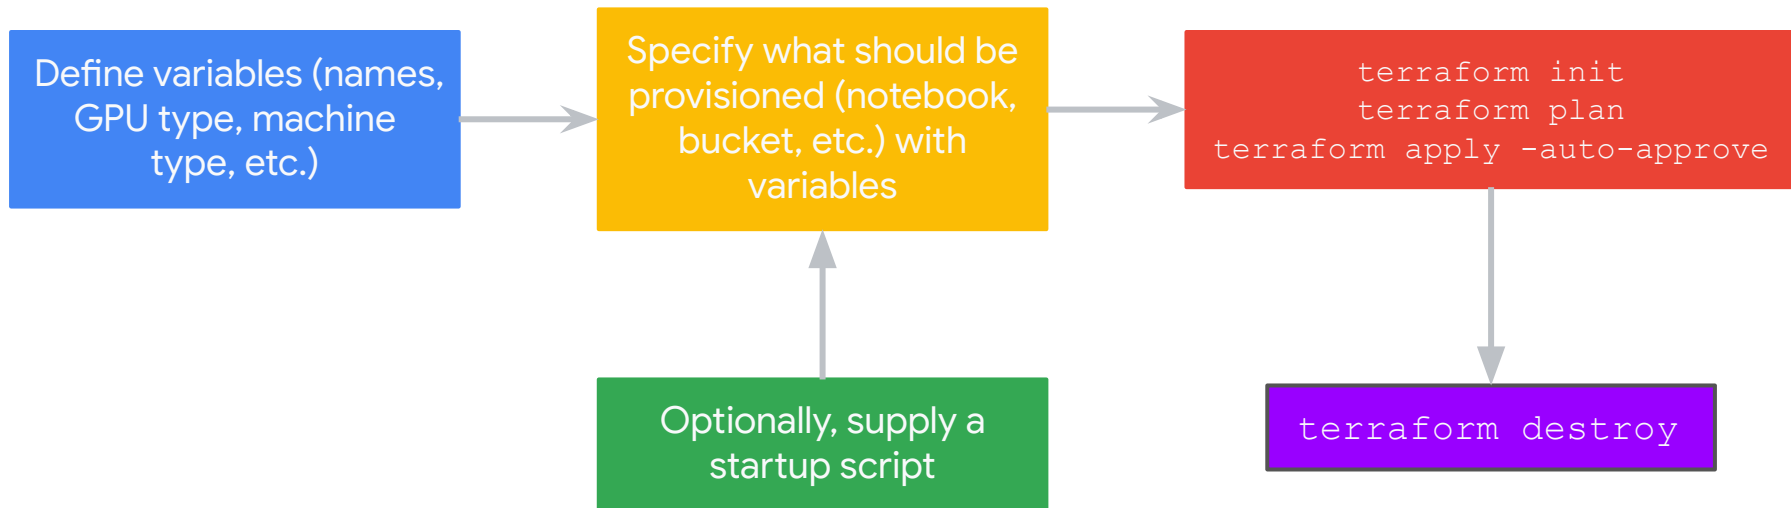
---





# Realize the same workflow but with code

---



# Standard anatomy for Terraform code

---

- `variables.tf` (define variables)
- `main.tf` (specify what should be provisioned with `variables.tf`)
- `scripts/`
  - `auto_shutdown.sh` (configure automatic shutdown)
  - `dependencies.sh` (configure additional dependencies)

`.tf` is the extension for HCL (Hashicorp Configuration Language)

# The `variables.tf` file

---

- GCP related information

- Project ID
- Location
- Region

- Notebook related information

- Name
- Machine type
- GPU type and count
- Image family

- GCS related information

- Name
- Region

# The variables.tf file

---

- GCP related information

```
variable "gcp" {  
  type = object(  
    {  
      project_id = string  
      location   = string  
      region     = string  
    }  
  )  
  default = {  
    project_id = "my-gcp-project"  
    location   = "us-central1-a"  
    region     = "us-central1"  
  }  
}
```

# The variables.tf file

---

- Notebook related information

```
variable "notebook" {  
  type = object(  
    {  
      notebook_name      = string  
      machine_type       = string  
      gpu_type           = string  
      ...  
    }  
  )  
  default = {  
    notebook_name      = "my-ml-nb"  
    machine_type       = "n1-standard-8"  
    gpu_type           = "NVIDIA_TESLA_T4"  
    ...  
  }  
}
```

# The variables.tf file

---

- GCS related information

```
variable "gcs" {  
  type = object(  
    {  
      name          = string  
      location      = string  
      force_destroy = bool  
    }  
  )  
  default = {  
    name          = "my-ml-bucket"  
    location      = "us-central1"  
    force_destroy = false  
  }  
}
```

# The `main.tf` file

---

- Provision notebook (Vertex AI Workbench) resource

```
resource "google_notebooks_instance" "notebook_instance" {  
  project      = var.gcp.project_id  
  name         = var.notebook.notebook_name  
  machine_type = var.notebook.machine_type  
  location     = var.gcp.location  
  ...  
}
```

# The `main.tf` file

---

- Provision notebook (Vertex AI Workbench) resource

```
resource "google_notebooks_instance" "notebook_instance" {  
  ...  
  network = ...  
  subnet  = ...  
  
  vm_image {  
    project      = local.image_project  
    image_family = var.notebook.image_family  
  }  
  
  accelerator_config {  
    type       = var.notebook.gpu_type  
    core_count = var.notebook.gpu_count  
  }  
  ...  
}
```



# The `main.tf` file

---

- Provision notebook (Vertex AI Workbench) resource

```
resource "google_notebooks_instance" "notebook_instance" {  
  ...  
  install_gpu_driver = var.notebook.install_gpu_driver  
  boot_disk_size_gb  = ...  
  data_disk_size_gb  = ...  
}
```

# The main.tf file

---

- Provision GCS bucket resource

```
resource "google_storage_bucket" "default" {  
  project          = var.gcp.project_id  
  name             = var.gcs.name  
  location        = var.gcs.location  
  storage_class    = "REGIONAL"  
  force_destroy    = var.gcs.force_destroy  
  uniform_bucket_level_access = true  
}
```

# Optionals

---

- Pass additional scripts to
  - Install additional dependencies
  - Automatically shutdown the instance w.r.t low CPU utilization
- Pass a URI of a predefined container image (`container_image`)
- Store the Terraform states to a remote location for better concurrency and collaboration

# State management: Shell-scripts vs. Terraform

---

- Terraform comes with state management
  - Keeps track of what's successfully provisioned or failed to provision
  - Control over the entire lifecycle of the infrastructure
  - Easy “planning” of the hardware to be provisioned for previewing changes
  - Enforces version control providing ease of collaboration

[bit.ly/next-22-tfm](https://bit.ly/next-22-tfm)

# Demo



# Wrapping up

---

- Manual hardware provisioning is faulty, introducing redundancy and rough edges.
- Streamlining the provisioning workflow establishes DRY and helps teams to have reproducibility when provisioning infrastructure.
- With infrastructure as code, we get:
  - Version control
  - Maintainability
  - Reusability

# Thank you

Sayak Paul (@RisingSayak)

[bit.ly/next-22-tfm](https://bit.ly/next-22-tfm)

Google Cloud

## Next '22

